

# HARSHA VARDHAN YELLELA

United States | +1-248-497-9965 | [harsha.yellela@gmail.com](mailto:harsha.yellela@gmail.com) | [har5ha.in](http://har5ha.in) | [LinkedIn](#) | [GitHub](#)

## SUMMARY

Software Engineer with expertise spanning **Machine Learning**, **Cloud Infrastructure**, and **Full-Stack Development**. Specialized in **LLM fine-tuning (QLoRA)**, **MLOps pipelines**, and **production Kubernetes deployments**. Built scalable systems including **94 AWS Lambda functions**, **8-microservice architectures**, and **enterprise Go backends**. Experienced in **ROS robotics**, **Graph Neural Networks**, and **Computer Vision**.

## EXPERIENCE

<b>Graduate Research Assistant – Agentic AI</b> <i>Lawrence Technological University</i>	<b>Jan 2025 – Dec 2025</b> <i>Southfield, MI</i>
<ul style="list-style-type: none"><li>Built and compared <b>no-code (n8n)</b> vs. <b>coded multi-agent systems (CrewAI + LangChain MCP)</b> for workflow automation and intelligent decision-making.</li><li>Deployed persistent <b>MCP agent services</b> on <b>AWS Fargate</b> and <b>Amazon EKS</b>, integrating <b>OpenSearch Serverless</b> for semantic search and RAG.</li><li>Designed hybrid pipelines combining <b>Bedrock-hosted models</b> with custom tools, achieving up to <b>70% reduction in manual process time</b>.</li></ul>	
<b>Infor India Pvt. Ltd.</b> <i>LN Technical Consultant</i>	<b>Apr 2022 – Dec 2023</b> <i>Hyderabad, India</i>
<ul style="list-style-type: none"><li>Developed modular, production-ready tools for global clients (<b>Ferrari, Boeing, Triumph</b>) by extending <b>Infor LN ERP</b> workflows.</li><li>Integrated <b>Infor ION</b> process flows with <b>AWS S3, Lambda, and API Gateway</b> for asynchronous file transfer and event-driven automation.</li><li><b>Containerized</b> business logic services using <b>Docker</b> and simulated Kubernetes-like orchestration with enterprise systems.</li></ul>	

## TECHNICAL SKILLS

- Languages:** Python, Go, TypeScript, JavaScript, SQL, Bash, Dart, C++
- ML/AI:** PyTorch, TensorFlow/Keras, PyTorch Geometric | LLM Fine-tuning (QLoRA, PEFT, TRL) | Transformers, HuggingFace
- LLM & Agents:** GPT-4, Gemini, Claude, Llama, Qwen | LangChain, LangGraph, CrewAI | RAG, Vector DBs (Qdrant, pgvector)
- Backend:** FastAPI, Flask, Gin (Go), Express.js | REST APIs, OpenAPI/Swagger, HATEOAS | JWT, OAuth
- Cloud (AWS):** Lambda, ECS Fargate, EKS, SageMaker, Bedrock | S3, DynamoDB, ECR | API Gateway, Kinesis, SNS
- DevOps:** Docker, Kubernetes | Terraform, CloudFormation, CDK | Jenkins, GitHub Actions | Prometheus, Grafana
- Robotics & Vision:** ROS (Robot Operating System), OpenCV, YOLO, CLIP | Point Cloud Processing, Sensor Fusion
- Frontend/Mobile:** React Native/Expo, Next.js, React | Flutter/Dart | TypeScript, CSS Modules
- Databases:** PostgreSQL, DynamoDB, MongoDB | Qdrant, pgvector (Vector) | Redis

## PROJECTS

<b>Blindspot – AI-Powered Cognitive Bias Detector</b> <i>Chrome Extension (Manifest V3), Claude API, JavaScript</i>   <i>Aurora Hackathon 2026</i>   <a href="#">GitHub</a>	<b>Jan 2026</b>
<ul style="list-style-type: none"><li>Built <b>real-time cognitive bias detector</b> Chrome extension that identifies decision-making flaws using <b>Claude API</b> for AI-powered analysis.</li><li>Implemented <b>proactive intervention system</b> monitoring high-risk sites (Amazon, eBay, DoorDash) detecting <b>6+ cognitive biases</b> including Sunk Cost Fallacy and Confirmation Bias.</li><li>Designed <b>context menu integration</b> and keyboard shortcuts for on-demand bias analysis with actionable reframe suggestions.</li></ul>	
<b>Resume Optimizer – QLoRA Fine-tuned LLM for ATS Optimization</b> <i>PyTorch, QLoRA, PEFT, TRL, Transformers, FastAPI, Ollama</i>   <a href="#">GitHub</a>	<b>Dec 2025</b>
<ul style="list-style-type: none"><li>Fine-tuned <b>Qwen3-4B</b> using <b>QLoRA (4-bit NF4 quantization)</b> with LoRA rank 16, alpha 32, reducing GPU memory to <b>18-22GB peak VRAM</b>.</li><li>Processed <b>1,800+</b> resumes to create <b>1,304 training examples</b> with structured JSON output, achieving <b>9.5/10 quality score</b> (GPT-5.1 evaluation).</li><li>Built <b>FastAPI REST API</b> for resume generation with <b>3-5 second inference time</b> on RTX 3090.</li></ul>	
<b>ML Sentiment Feedback Loop – Production MLOps Microservices</b> <i>AWS (ECS Fargate, SageMaker, S3), Terraform, GitHub Actions, Docker</i>   <a href="#">GitHub</a>	<b>Dec 2025</b>
<ul style="list-style-type: none"><li>Built <b>8-microservice architecture</b> (API Gateway, Inference, Feedback, Model Registry, Evaluation, Retraining, Notification, Model Init) with <b>independent scaling</b>.</li><li>Implemented <b>complete ML feedback loop</b> with <b>auto-retraining</b>, model versioning, and <b>SageMaker integration</b> for training jobs.</li><li>Configured <b>GitHub Actions CI/CD</b> with <b>Terraform IaC</b> for automated infrastructure provisioning and container deployments.</li></ul>	

## FieldFuze Backend – Enterprise Go REST API

Dec 2024

*Go (Gin), AWS DynamoDB, JWT, Docker, GitHub Actions*

- Built **production-ready REST API** in **Go 1.23** with **multi-tenant RBAC architecture** and **permission inheritance**.
- Designed **modular architecture** (controller → service → repository pattern) with **comprehensive middleware** (JWT auth, CORS, logging).
- Implemented **infrastructure automation worker** for DynamoDB table/index management with **extensive unit tests** for all layers.

## Stretch2 Robot – Autonomous Navigation & Grasping

Dec 2024

*ROS, Python, OpenCV, YOLO, CLIP, Point Cloud Processing | GitHub*

- Built **autonomous object cluster detection** using **ROS services** with **multi-strategy arm manipulation** (mean/max/random fallback).
- Integrated **YOLO object detection** with **CLIP segmentation** for multi-modal understanding on **Stretch2 mobile manipulator**.
- Developed **point cloud processing** with farthest point sampling and **location change detection** to avoid redundant scans.

## Lambda Microservices Platform – Enterprise Serverless Backend

Dec 2024

*AWS Lambda, DynamoDB, API Gateway, Terraform, CloudFormation | Stripe, Twilio, DocuSign, QuickBooks*

- Developed **94 AWS Lambda functions** for complete SaaS platform with **Terraform** and **CloudFormation** infrastructure.
- Integrated **10+ third-party services**: Stripe Connect (payments), DocuSign (e-signatures), Twilio (communications), QuickBooks (accounting), EagleView (aerial imagery).
- Built **automated deployment pipeline** with Git change detection, deploying only modified functions to reduce deployment time.

## EDUCATION

### Lawrence Technological University

Jan 2024 – Dec 2025

*Master of Science in Computer Science · GPA: 3.6/4.0*

Southfield, MI

- Relevant Coursework: Machine Learning, Artificial Intelligence, Natural Language Processing, Intelligent Robotics (ROS), Agentic AI Research

### Geethanjali College of Engineering & Technology

Aug 2018 – Aug 2022

*Bachelor of Technology in Computer Science & Engineering · GPA: 7.5/10 (~3.0/4.0)*

Hyderabad, Telangana

- Relevant Coursework: Deep Learning & Python, Machine Learning Foundations, Software Engineering, Internet of Things

## ACHIEVEMENTS

- Selected for Amazon Nova AI Challenge: Trusted AI Track (2025)
- Participated in RSNA Pneumonia Detection Challenge; ranked in upper quartile with VGG19 transfer learning model (2024)
- Built production SaaS platform with 94 Lambda functions serving real customers for field service management (2024)
- Gold Medalist in Indian National Mathematical Olympiad (INMO) (2012)